

An Introduction to the AMD Optimizing C/C++ Compiler (AOCC) and the AMD Optimizing CPU Libraries (AOCL)

Marty Kandes, Ph.D.

Computational & Data Science Research Specialist
High-Performance Computing User Services Group
San Diego Supercomputer Center
University of California, San Diego

XSEDE Webinar
Wednesday, April 21st, 2021
09:45 AM - 10:35 AM PST

XSEDE Code of Conduct

Code of Conduct

This external code of conduct for XSEDE-sponsored events represents XSEDE's commitment to providing an inclusive and harassment-free environment in all interactions regardless of race, age, ethnicity, national origin, language, gender, gender identity, sexual orientation, disability, physical appearance, political views, military service, health status, or religion. The code of conduct below extends to all XSEDE-sponsored events, services, and interactions.

XSEDE is committed to providing an inclusive and harassment-free environment in all interactions regardless of race, age, ethnicity, national origin, language, gender, gender identity, sexual orientation, disability, physical appearance, political views, military service, health status, or religion. This commitment extends to all XSEDE-sponsored events and services (in-person training, webinars, committee meetings, networking functions, online forums, chat rooms, and social media) and any interaction including staff-to-participant, participant-to-participant, and participant-to-staff. Participants are all individuals who are not XSEDE staff who attend and participate including but not limited to administrators, faculty, students, researchers, and research computing professionals. As a project that aims to share ideas and freedom of thought and expression, it is essential that the interaction between participants, users of XSEDE services, and XSEDE staff take place in an environment that recognizes the inherent worth of every person by being respectful of all. The XSEDE project does not tolerate harassment in any form. Harassment is any form of behavior intended to exclude, intimidate, or cause discomfort. Harassment includes, but is not limited to, the use of abusive or degrading language, intimidation, staking, harassing photography or recording, inappropriate physical contact, and unwelcome sexual attention. All XSEDE users, participants, and staff are governed by local laws and their organization's code of conduct and policies.

Anyone who experiences, observes, or has knowledge of threatening behavior are expected to immediately report the incident to a member of the event organizing committee, XSEDE staff, or one of the XSEDE Ombudspersons listed below, or by using the [online form](#). XSEDE reserves the right to take appropriate action.

XSEDE ombudspersons:

- Linda Akl, Southeastern Universities Research Association (lakl@luru.org)
- Lizzanne Destefano, Georgia Tech (lizzanne.destefano@ccsmc.gatech.edu)
- Ken Hackworth, Pittsburgh Supercomputing Center (hackworth@psc.edu)
- Bryan Snead, Texas Advanced Computing Center (bsnead@tacc.utexas.edu)

Key Points

- XSEDE provides an inclusive and harassment-free environment
- Code of conduct applies to all XSEDE-sponsored events and services

Related Links

[Anonymous Reporting Form](#)

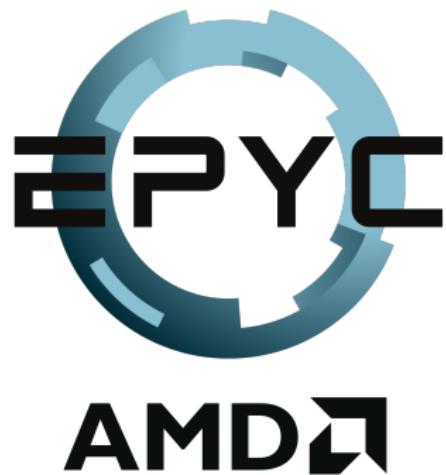
Contact Information

Linda Akl

<https://www.xsede.org/codeofconduct>

Today

- ▶ AMD Optimizing C/C++ Compiler
 - Compiler Options: gpecn1Dtdpbc
- ▶ AMD Optimizing CPU Libraries
 - AMD Math Library: savage
 - AMD BLIS: pave; dgemm
- ▶ Summary & References
- ▶ Q & A



AMD

AMD Optimizing C/C++ Compiler (AOCC)

AOCC is a high performance, optimized set of LLVM-based compilers. They provide developers with all of the essential tools to build and optimize C, C++, and Fortran applications targeting AMD processors. Some features include:

- ▶ Enabled, tuned and optimized for AMD EPYC architectures
- ▶ Based on LLVM Compiler Infrastructure
- ▶ Flang is default Fortran front-end with F2008, Real 128 features
- ▶ OpenMP support
- ▶ RHEL 8, CentOS 8, SLES 15, and Ubuntu 20.04 LTS
- ▶ Spack support

<https://developer.amd.com/amd-aocc>

LLVM

The LLVM Project is a collection of modular and reusable compiler and toolchain technologies. Its core libraries provide a modern source- and target-independent optimizer, along with code generation support for many popular CPUs architectures. These libraries are built around a well specified code representation known as the LLVM intermediate representation ("LLVM IR").

<https://llvm.org>

Clang

The Clang project provides a language front-end and tooling infrastructure for languages in the C language family for the LLVM project. It is considered to be a production quality C, Objective-C, C++ and Objective-C++ compiler when targeting X86-32, X86-64, and ARM architectures. Clang supports C++11, C++14 and C++17.

<https://clang.llvm.org>

Flang

Flang (also known as "Classic Flang") is a Fortran compiler targeting LLVM. It is an open-sourced version of pgfortran, a commercial Fortran compiler from PGI/NVIDIA. It is different from the new Flang (formerly known as "F18"; see <https://flang.llvm.org>), which has been part of the LLVM project since 2020, although both are developed by the same community. It is also unrelated to other projects of the same name, such as <https://github.com/llvm-flang/flang> and <https://github.com/isanbard/flang>.

Classic Flang is used in several downstream commercial projects, and continues to be maintained, but the plan is to replace Classic Flang with the new Flang in the future.

<https://github.com/flang-compiler/flang>

AOCC Compiler Options

AOCC compiler (with Flang - Fortran Front-End)

Latest release: 2.3, Dec 2020

<https://developer.amd.com/amd-aocc/>

Architecture		Other options		
Generate instructions that run on 2nd Gen EPYC/ RYZEN	-march=znver2	Enable faster, less precise math operations	-ffast-math -freciprocal-math	
Generate instructions for the local machine	-march=native	OpenMP® threads and affinity (N number of cores)	export OMP_NUM_THREADS=N export GOMP_CPU_AFFINITY="0-[N-1]"	
Optimization Levels			Enabling Vector library	
Disable all optimizations	-O0	Link to Vector library	-vector-library=LIBMVEC	
Minimal level speed and code optimization	-O1	Link to AMD library	-L/libm-install-dir/lib -lvec -L/libm-install-dir/lib -lamdlibm	
Moderate level optimization (default)	-O2 / -O	For Fortran workloads		
Aggressive optimizations	-O3	Compile free form FORTRAN	-ffree-form	
Maximize performance	-Ofast	Enable precise math operations in FORTRAN	-kieee	
Enable Link Time Optimization	-fto	AMD Optimized Libraries		
Enable unrolling	-funroll-loops	Latest release: 2.2, Jun 2020		
Enable aggressive loop optimizations	-enable-loop-versioning-lcm -enable-partial-unswitch -unroll-aggressive	https://developer.amd.com/amd-aocl/		
Enable aggressive inline optimizations	-function-specialize -finline-aggressive	AMD uProf (Performance & Power Profiler)		
Enable aggressive vectorization	-enable-strided-vectorization -enable-epilog-vectorization	Latest release: 3.3, Jul 2020		
Enable memory layout optimizations	-fremap-arrays (use with -fto)	https://developer.amd.com/amd-uprof/		
Profile Guided optimizations	-profile-instr-generate (1st invocation) -profile-instr-use (2nd invocation)	Other options		
OpenMP®	-fopenmp	Enable generation of code that follows IEEE arithmetic	-mieee-fp	
For enabling streaming stores, memory bandwidth workloads	-fnt-store	Enable faster, less precise math operations	-ffast-math	
Enable removal of un-used array computation	-reduce-array-computation=3	Compile free form FORTRAN	-ffree-form	

GNU compiler collection (gcc, g++, gfortran)

Latest release: 10.3, Jul 2020

Recommended version : 9.3

<http://gcc.gnu.org>

Architecture	
Generate instructions that run on 2nd Gen EPYC/ RYZEN	-march=znver2
Generate instructions for the local machine	-march=native
Optimization Levels	
Disable all optimizations (default)	-O0
Minimal level speed and code optimizations	-O1 / -O
Moderate level optimizations	-O2
Aggressive optimizations	-O3
Maximize performance	-Ofast
Additional Optimizations	
Link time optimization	-ftlo
Enable unrolling	-funroll-all-loops
Generate memory preload instructions	-prefetch-loop-arrays --param prefetch-latency=300
Profile-guided optimization	-profile-generate (1st invocation) -fprofile-use (2nd invocation)
OpenMp®	-fopenmp
Other options	
Enable generation of code that follows IEEE arithmetic	-mieee-fp
Enable faster, less precise math operations	-ffast-math
Compile free form FORTRAN	-ffree-form
OpenMp® threads and affinity (N number of cores)	export OMP_NUM_THREADS=N export GOMP_CPU_AFFINITY="0-[N-1]"
Link to AMD library	-L/libm-install-dir/lib -lamdlibm

Example: A Real-World Science Code's First Test Drive of AOCC

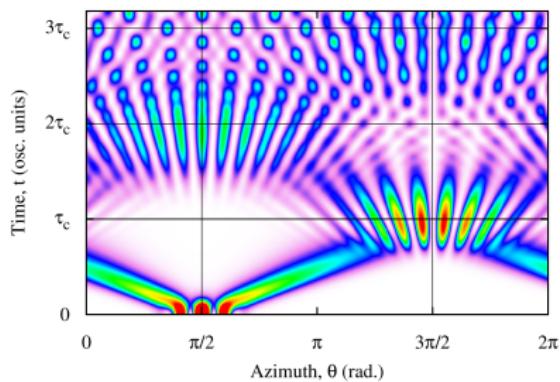
gpecn1Dtdpbc: Overview

- ▶ Gross-Pitaevskii Equation (GPE)
- ▶ Crank-Nicolson (CN) scheme;
2nd-order central differences
- ▶ One-Dimensional (1D) ring topology
- ▶ Time-Dependent (TD)
- ▶ Periodic Boundary Conditions (PBC); Sherman-Morrison algorithm

```
PROGRAM GPECN1DTPBC
IMPLICIT NONE

Parameter Declarations:
REAL, PARAMETER :: PI = 3.1415926535897932384626433832795828842E0
REAL :: RADIUS
REAL :: ROTATION_RATE
REAL :: GAMMA_FACTOR
COMPLEX, ALLOCATABLE, DIMENSION(:) :: WAVEFUNCTION
COMPLEX, ALLOCATABLE, DIMENSION(:) :: EXTERNAL_POTENTIAL
COMPLEX, ALLOCATABLE, DIMENSION(:) :: NONLINEAR_COUPLING

Input Variable and Array Declarations:
INTEGER :: NUMBER_OF_GRID_POINTS
INTEGER :: NUMBER_OF_TIME_STEPS
INTEGER :: NUMBER_OF_TIME_STEPS_BEFORE_WRITE
INTEGER :: NUMBER_OF_ITERATIONS_PER_TIME_STEP
INTEGER :: IMAGINARY_TIME_SWITCH
```



[https://github.com/mkandes/
gpecn1Dtdpbc](https://github.com/mkandes/gpecn1Dtdpbc)

$$i \frac{\partial}{\partial t} \psi(\theta, t) = \left[-\frac{1}{2R^2} \frac{\partial^2}{\partial \theta^2} - \Omega L_z + V(\theta, t) + g_s |\psi(\theta, t)|^2 \right] \psi(\theta, t)$$

gpecn1Dtdpbc: Computational Features & Operations

- ▶ Fortran 90
- ▶ Serial, single-core
- ▶ Double precision, complex arrays
- ▶ Element-wise array operations
- ▶ Intrinsics:
CMPLX, CONJG,
SUM
- ▶ LAPACK: ZGTSV

```
20 #     Specify general user-defined compilation options.
21
22 COMPILER    := gfortran
23 OPTIMIZATION := ON
24 OPENMP      := OFF
25 DEBUG       := OFF
26 PROFILE     := OFF
27
28 #     Set compiler-specific options.
29
30 ifeq ($(COMPILER),gfortran)
31
32 STANDARD_OPTIONS   := -fimplicit-none -fmodule-private \
33                         -ffree-form -ffree-line-length-none -std=gnu
34 INTEGER_OPTIONS    := -fdefault-integer-8
35 REAL_OPTIONS       := -fdefault-real-8 -fdefault-double-8
36 OPTIMIZATION_OPTIONS := -O3 -ffast-math -mtune=native
37 OPENMP_OPTIONS     := -fopenmp
38 CHECK_OPTIONS      := -fcheck=all
39 DEBUG_OPTIONS      := -ffpe-trap=invalid,overflow -fbacktrace \
40                         -fdump-core -finit-real=nan
41 WARNING_OPTIONS    := -Wall -fmax-errors=0 -Wno-array-temporaries \
42                         -Warray-bounds -Wcharacter-truncation \
43                         -Wline-truncation -Wconversion-extra \
44                         -Wimplicit-interface -Wimplicit-procedure \
45                         -Wunderflow -Wextra -Wuninitialized
46 PROFILE_OPTIONS    := -pg
47
48 endif
49
50 ifeq ($(COMPILER),ifort)
51
52 STANDARD_OPTIONS   := -implicitnone -free -stand none -module build
53 INTEGER_OPTIONS    := -integer-size 64
54 REAL_OPTIONS       := -real-size 64 -double-size 64
55 OPTIMIZATION_OPTIONS := -O3 -fast -mtune=native
56 OPENMP_OPTIONS     := -fopenmp
57 CHECK_OPTIONS      := -check all
58 DEBUG_OPTIONS      := -debug all
59 WARNING_OPTIONS    := -warn all
60 PROFILE_OPTIONS    := -pg
61
62 endif
63
```

gpecn1dtdpbc: Runtime Performance

ARCH	COMPILER	MIN (s)	MAX (s)	AVG (s)	STDEV (s)
EPYC 7742	gfortran 8.3.1	897.5	1001.3	949.0	52.9
	gfortran 10.2.0	1028.4	1098.0	1080.1	27.4
	ifort 19.1.1.217	1445.8	1669.1	1529.1	76.9
	aocc/flang 2.2.0	1797.7	1980.0	1907.6	85.3
Xeon Gold 6248	ifort 19.0.5.281	795.5	873.8	810.2	22.8

```
gfortran -fimplicit-none -fmodule-private -ffree-form -ffree-line-length-none -std=gnu -fdefault-integer-8 -fdefault-real-8 -fdefault-double-8 -O3 -ffast-math -mtune=native

ifort -fimplicitnone -free -stand none -module build -integer-size 64 -real-size 64 -double-size 64 -O3 -fast -mtune=native

flang -Mfreeform -Mstandard -fdefault-integer-8 -fdefault-real-8 -O3 -ffast-math -mtune=native
```

gpecn1dtdpbc: Shared (L3) Cache Contention

[mkandes@exp-2-02 ~]\$ top

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+ COMMAND
106813	zhanglab	28	0	985924	256072	1260	99.7	0.1	0.1	0:11.38 cas_569986_2A_
109370	zhanglab	28	0	985924	256072	1260	99.7	0.1	0.1	0:14.06 cas_569986_3A_
138249	unctrl	28	0	31216	16536	7996	99.7	0.1	0.1	0:16.34 python
138444	elegraf	28	0	118140	118140	42748	S	10.3	0.0	0:00:32.14 telegraf
88889	zhanglab	28	0	1856448	659344	4168	S	3.6	0.0	0:02.54 zabbix
111433	zhanglab	28	0	79512	26989	3656	S	2.0	0.0	0:00:48 scontrol
178001	nlcd	28	0	456856	6616	4168	S	1.7	0.0	0:05.67 nlcd
113807	nikandes	28	0	67996	3723	3929	R	1.0	0.0	0:00:11 top
78848	robpearc	28	0	58936	27672	4568	S	0.7	0.0	0:08.92 perl
40881	root	28	0	0	0	0	I	0.3	0.0	0:00:18 kworker/30:2-mm
4196	robpearc	28	0	58936	27672	4568	S	0.3	0.0	0:18.88 perl
4958	robpearc	28	0	58948	27676	4568	S	0.3	0.0	0:28.88 perl
7876	robpearc	28	0	58948	27676	4568	S	0.3	0.0	0:48.41 perl
18474	root	28	0	210876	33468	31344	S	0.3	0.0	0:15.43 gssd_nss
13678	robpearc	28	0	58948	27676	4568	S	0.3	0.0	0:35.03 perl
17919	root	28	0	0	0	0	S	0.3	0.0	26:12.17 kiblnd_sd_01_03
17920	root	28	0	0	0	0	S	0.3	0.0	2:07.97 kiblnd_sd_02_00
18137	root	28	0	0	0	0	S	0.3	0.0	0:42.71 ptlrpcd_02_05
18423	root	28	0	3900788	772424	106932	S	0.3	0.3	164:41.68 cmd
21447	robpearc	28	0	58936	27672	4568	S	0.3	0.0	0:19.58 perl
22225	robpearc	28	0	58948	27676	4568	S	0.3	0.0	0:52.46 perl
29138	robpearc	28	0	58944	27680	4568	S	0.3	0.0	1:16.16 perl
42210	robpearc	28	0	58944	27680	4568	S	0.3	0.0	0:13.92 perl
51239	robpearc	28	0	58948	27676	4568	I	0.3	0.0	0:37.51 perl
81616	robpearc	28	0	58948	27676	4568	S	0.3	0.0	0:03.83 perl
84881	robpearc	28	0	58944	27680	4568	S	0.3	0.0	0:38.35 perl
88289	robpearc	28	0	58936	27672	4568	I	0.3	0.0	0:32.31 perl
96322	robpearc	28	0	58948	27676	4568	S	0.3	0.0	0:28.54 perl
117982	robpearc	28	0	58948	27676	4568	S	0.3	0.0	0:47.26 perl
129892	root	28	0	279936	7876	5088	S	0.3	0.0	0:01.18 slurmstepd
1 root	root	28	0	164348	12880	4436	S	0.4	0.0	14:59.20 systemd
2 root	root	28	0	0	0	0	S	0.0	0.0	0:18.94 kernel
3 root	root	28	0	0	0	0	S	0.0	0.0	0:00:00 rcu_ksm
4 root	root	28	0	0	0	0	I	0.0	0.0	0:00:00 rcu_ksm
6 root	root	28	0	0	0	0	I	0.0	0.0	0:00:00 kworker/0:0-H:kbl
9 root	root	28	0	0	0	0	I	0.0	0.0	0:00:00 percpu_wq
10 root	root	28	0	0	0	0	S	0.0	0.0	0:49.88 ksaftrid/q
11 root	root	28	0	0	0	0	I	0.0	0.0	46:33.17 rcu_sched
12 root	rt	28	0	0	0	0	S	0.0	0.0	0:00:04 migration@
13 root	rt	28	0	0	0	0	S	0.0	0.0	0:00:04 watchdog@
14 root	root	28	0	0	0	0	S	0.0	0.0	0:00:00 cpuhwp/
15 root	root	28	0	0	0	0	S	0.0	0.0	0:00:00 cpuhwp/1
16 root	rt	28	0	0	0	0	S	0.0	0.0	0:00:24 watchdog/1

1798.7 s

[mkandes@exp-2-14 ~]\$ top

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+ COMMAND
103137	huantran	28	0	6647796	358348	169876	R	161.7	6.1	3:51.95 vasp_gam
11796	misbaazn	28	0	1922884	1.1g	22628	R	166.6	6.4	1:47.23 vasp_std
11799	misbaazn	28	0	1922836	1.1g	32888	R	166.0	6.4	1:55.12 vasp_std
11800	misbaazn	28	0	1922644	1.1g	32448	R	166.0	6.4	1:05.28 vasp_std
11801	misbaazn	28	0	1922648	1.1g	32492	R	166.0	6.4	1:44.12 vasp_std
11803	misbaazn	28	0	1922624	1.1g	32648	R	166.0	6.4	1:37.59 vasp_std
11805	misbaazn	28	0	1922888	1.1g	32544	R	166.0	6.4	1:42.16 vasp_std
11811	misbaazn	28	0	1922468	1.1g	32924	R	166.0	6.4	1:48.21 vasp_std
11812	misbaazn	28	0	1922488	1.1g	32920	R	166.0	6.4	1:39.12 vasp_std
11815	misbaazn	28	0	1922344	1.1g	32464	R	166.0	6.4	1:44.15 vasp_std
11817	misbaazn	28	0	1922516	1.1g	32328	R	166.0	6.4	1:44.23 vasp_std
11818	misbaazn	28	0	1922972	1.1g	32356	R	166.0	6.4	1:43.21 vasp_std
11819	misbaazn	28	0	1922428	1.1g	32484	R	166.0	6.4	1:44.86 vasp_std
11820	misbaazn	28	0	1922716	1.1g	32324	R	166.0	6.4	1:44.82 vasp_std
11822	misbaazn	28	0	1922444	1.1g	32340	R	166.0	6.4	1:45.15 vasp_std
11824	misbaazn	28	0	1922448	1.1g	32524	R	166.0	6.4	1:43.57 vasp_std
56471	umai	28	0	31248	16164	8012	R	166.0	6.0	6:57.37 python
56599	umai	28	0	31092	16432	8032	R	166.0	6.0	6:57.32 python
56649	umai	28	0	31224	16176	8032	R	166.0	6.0	6:51.35 python
56735	umai	28	0	31200	16168	8032	R	166.0	6.0	6:52.31 python
56847	umai	28	0	31228	16192	8032	R	166.0	6.0	6:51.53 python
56869	umai	28	0	31212	16156	8032	R	166.0	6.0	6:57.39 python
103122	huantran	28	0	6654508	395846	192184	R	166.0	6.2	3:46.66 vasp_gam
103124	huantran	28	0	6652316	399992	198812	R	166.0	6.2	3:52.32 vasp_gam
103125	huantran	28	0	6666664	401128	197884	R	166.0	6.2	3:52.35 vasp_gam
103126	huantran	28	0	6666782	381492	179848	R	166.0	6.1	3:51.27 vasp_gam
103127	huantran	28	0	6655598	375232	176486	R	166.0	6.1	3:51.44 vasp_gam
103128	huantran	28	0	6666852	376782	179184	R	166.0	6.1	3:51.38 vasp_gam
103131	huantran	28	0	6657950	359788	162440	R	166.0	6.1	3:51.32 vasp_gam
103133	huantran	28	0	6644976	356782	159289	R	166.0	6.1	3:51.34 vasp_gam
103134	huantran	28	0	6658456	358748	161688	R	166.0	6.1	3:51.09 vasp_gam
103135	huantran	28	0	6664202	359697	159289	R	166.0	6.0	3:51.09 vasp_gam
103136	huantran	28	0	6664432	360044	151616	R	166.0	6.1	3:51.97 vasp_gam
118897	zhanglab	28	0	98524	266652	1260	R	166.0	6.1	1:16.37 29 cas_569919_4F
128072	zhanglab	28	0	98524	217724	1260	R	166.0	6.1	748:59.35 cas_568916_5F
11792	misbaazn	28	0	1948800	31	31772	R	98.3	6.0	1:49.43 vasp_std
11798	misbaazn	28	0	1922576	1.1g	32732	R	98.3	6.4	1:44.85 vasp_std
11802	misbaazn	28	0	1922756	1.1g	32532	R	98.3	6.4	1:44.40 vasp_std
11804	misbaazn	28	0	1922864	1.1g	32584	R	98.3	6.4	1:44.24 vasp_std
11806	misbaazn	28	0	1922572	1.1g	32520	R	98.3	6.4	1:37.40 vasp_std
11807	misbaazn	28	0	1922828	1.1g	32712	R	98.3	6.4	1:44.11 vasp_std
11808	misbaazn	28	0	1922592	1.1g	32264	R	98.3	6.4	1:38:45 vasp_std



gpecn1Dtdpbc: Compiler Options & Build Warnings

```
flang -Mfreeform -Mstandard -fdefault-integer-8 -fdefault-real-8 -O3 -ffast-math
      -mtune=native -c source/gpecn1Dtdpbc.f -o build/gpecn1Dtdpbc.o
F90-W-0155-The type of FLOAT is now double precision with -r8 (source/
               gpecn1Dtdpbc.f: 166)
    0 inform, 1 warnings, 0 severes, 0 fatal for gpecn1Dtdpbc
flang -Mfreeform -Mstandard -fdefault-integer-8 -fdefault-real-8 -O3 -ffast-math
      -mtune=native -c libraries/lapack/zgtsv.f -o build/zgtsv.o
5 F90-W-0173-PGI Fortran extension: nonstandard use of data type length specifier
    (libraries/lapack/zgtsv.f: 11)
F90-W-0173-PGI Fortran extension: nonstandard use of data type length specifier
    (libraries/lapack/zgtsv.f: 71)
F90-W-0173-PGI Fortran extension: nonstandard use of data type length specifier
    (libraries/lapack/zgtsv.f: 76)
    0 inform, 3 warnings, 0 severes, 0 fatal for zgtsv
flang -Mfreeform -Mstandard -fdefault-integer-8 -fdefault-real-8 -O3 -ffast-math
      -mtune=native -c libraries/lapack/xerbla.f -o build/xerbla.o
10 flang -Mfreeform -Mstandard -fdefault-integer-8 -fdefault-real-8 -O3 -ffast-math
      -mtune=native build/gpecn1Dtdpbc.o build/zgtsv.o build/xerbla.o -o
           gpecn1Dtdpbc.x
clang-10: warning: argument unused during compilation: '-Mfreeform' [-Wunused-
           command-line-argument]
clang-10: warning: argument unused during compilation: '-Mstandard' [-Wunused-
           command-line-argument]
clang-10: warning: argument unused during compilation: '-fdefault-integer-8' [-
          Wunused-command-line-argument]
clang-10: warning: argument unused during compilation: '-fdefault-real-8' [-
          Wunused-command-line-argument]
```

AMD Optimizing CPU Libraries (AOCL)

AOCL is a set of numerical libraries tuned specifically for AMD EPYC processor family. They have a simple interface to take advantage of the latest hardware innovations. The tuned implementations of industry standard math libraries enable fast development of scientific and high performance computing projects.

- ▶ AMD Math Library (LibM)
- ▶ AMD BLIS, libFLAME, ScaLAPACK, AOCL-Sparse
- ▶ AMD FFTW
- ▶ AMD Random Number Generator and Secure RNG Library

<https://developer.amd.com/amd-aocl>

AMD Math Library (LibM)

AMD LibM is a software library containing a collection of basic math functions optimized for x86-64 processor based machines. It provides many routines from the list of standard C99 math functions. AMD LibM is a C library, which users can link in to their applications to replace compiler-provided math functions. Generally, programmers access basic math functions through their compiler, but those who want better accuracy or performance than their compiler's math functions can use this library to help improve their applications.

<https://developer.amd.com/amd-aocl/amd-math-library-libm>

<https://github.com/amd/aocl-libm-ose>

AMD BLIS (BLAS) Library

BLIS is a portable software framework for instantiating high-performance BLAS-like dense linear algebra libraries. The framework was designed to isolate essential kernels of computation that, when optimized, enable optimized implementations of most of its commonly used and computationally intensive operations. Select kernels have been optimized for the AMD EPYC processor family. The optimizations are done for single and double precision routines.

<https://developer.amd.com/amd-aocl/blas-library>

<https://github.com/amd/blis>

AMD libFLAME

libFLAME is a portable library for dense matrix computations, compatible with Netlib LAPACK specification. It includes a compatibility layer, FLAPACK, which includes complete LAPACK implementation. The library provides scientific and numerical computing communities with a modern, high-performance dense linear algebra library that is extensible, easy to use, and available under an open source license.

libFLAME is a C-only implementation and does not depend on any external FORTRAN libraries including LAPACK. There is an optional backward compatibility layer, lapack2flame that maps LAPACK routine invocations to their corresponding native C implementations in libFLAME. This allows legacy applications to start taking advantage of libFLAME with virtually no changes to their source code.

<https://github.com/amd/libflame>

AMD ScaLAPACK

ScaLAPACK is a library of high-performance linear algebra routines for parallel distributed memory machines. It depends on external libraries including BLAS and LAPACK for Linear Algebra computations. AMD's optimized version of ScaLAPACK enables using BLIS and libFLAME library that have optimized dense matrix functions and solvers for AMD EPYC processor family CPUs.

<https://developer.amd.com/amd-aocl/scalapack>

<https://github.com/amd/scalapack>

AMD AOCL-Sparse

AOCL-Sparse is a library that contains basic linear algebra subroutines for sparse matrices and vectors optimized for AMD EPYC family of processors. It is designed to be used with C and C++. Current functionality of sparse library supports SPMV function with CSR and ELLPACK formats.

<https://developer.amd.com/amd-aocl/aocl-sparse>

<https://github.com/amd/aocl-sparse>

AMD FFTW

FFTW is a comprehensive collection of fast C routines for computing the Discrete Fourier Transform (DFT) and various special cases thereof. It is an open-source implementation of the Fast Fourier transform algorithm. It can compute transforms of real and complex-values arrays of arbitrary size and dimension. AMD optimized FFTW includes selective kernels and routines optimized for the AMD EPYC processor family.

<https://developer.amd.com/amd-aocl/fftw/>

<https://github.com/amd/amd-fftw>

AMD Random Number Generator Libraries

AMD Random Number Generator (RNG) Library is a pseudorandom number generator library. It provides a comprehensive set of statistical distribution functions which are founded on various underlying uniform distribution generators (base generators) including Wichmann-Hill and Mersenne Twister. The library contains five base generators and twenty-three distribution generators. Users can also supply a custom built generator as the base generator for all of the distribution generators.

AMD Secure RNG is a library that provides APIs to access the cryptographically secure random numbers generated by AMD's hardware-based random number generator implementation.

<https://developer.amd.com/amd-aocl/rng-library>

Example: Trancendental Functions & AMD Math Library (LibM)

Savage Benchmark

16-Bit Toolbox (Text begins on page 120)

Listing One BASIC and PL-1 versions of Bill Savage's accuracy and speed test for high level languages.

```
100  ' Time and General Accuracy Test Program
110  DEFINT I
120  ILOOP=2500
130  A=1
140  FOR I=1 TO ILOOP-1
150    A = TAN(ATN(EXP(LOG(SQR(A*A)))) + 1
160  NEXT I
170  PRINT USING "A#####.#####";A
180  STOP
```

"An interesting program to test the speed and accuracy of various high-level languages has been sent to us by Bill Savage of Microfloat in Houston, Texas ... The final result of each test should be 2500.000; any variations from this figure reflect errors in the language's floating-point algorithms or rounding methods. The arctangent in the central function is designed to accentuate such discrepancies and make them more visible."

- R. Duncan, "16-bit Software Toolbox", Dr. Dobb's Journal, Number 83, September 1983, p. 120

Savage Benchmark: Runtime Performance (circa 1983)

With hardware floating point libraries from MicroFloat:

Language	Version	Result	Time (sec.)
PL/I-86 with 8087	1.01	2477.244	3.7
8080 Assembler with 8232	RMAC	2499.955	10.2
PL/I-80 with 8232	1.40	2499.995	10.4
BASIC-80 with 8232	5.20	2500.000	10.7
Fortran-80 with 8232	3.40	2499.995	12.5

With standard floating point libraries:

Language	Version	Result	Time (sec.)
BASIC-86 Interpreter	5.20	2179.850	92.2
Fortran-80	3.40	2304.863	140.8
BASIC-80 Compiler	5.20	2304.860	140.8
BASIC-80 Interpreter	5.20	2304.860	174.9
PL/I-86	1.01	1641.758	179.6
PL/I-80	1.30	1641.758	254.4

Table 1.

Comparisons of execution speed and accuracy for commonly available high-level languages, with and without hardware floating point support. The 8087 and 8232 support libraries are marketed under the names FLOAT87 and FLOAT80, which are trademarks of MicroFloat.

Savage Benchmark: Runtime Performance (circa 2021)

Measured on Expanse's AMD EPYC 7742 compute nodes when linking against the standard libm distributed with CentOS 8.

n (M)	a(n)	gcc 10.2.0	icc 19.1.1.217	aocc/clang 2.2.0
1	999999.997712	0.118	0.071	0.116
10	9999816.516784	1.112	0.675	1.161
20	19996559.553870	520.955	1.338	581.503
50	49671498.766991		3.343	
100	94906263.696971		6.66	
1000	94906263.696971		66.61	

```
gcc -O3 -mtune=native savage.c -o savagec.x -lm  
icc -O3 -mtune=native savage.c -o savagec.x -lm  
clang -O3 -mtune=native savage.c -o savagec.x -lm
```

<https://github.com/mkandes/savage>

Savage Benchmark: Runtime Performance w/AMD LibM

Measured on Expanse's AMD EPYC 7742 compute nodes when linking against the AMD Math Library (LibM).

n (M)	a(n)	gcc 10.2.0	icc 19.1.1.217	aocc/clang 2.2.0
1	999999.997712	0.105	0.096	0.116
10	9999816.516784	1.024	0.945	1.047
20	19996559.553870	2.045	1.851	1.891
50	49671498.766991	5.102	4.723	4.547
100	94906263.696971	10.221	9.417	10.1
1000	94906263.696971	102.167	94.342	100.068

```
declare -xr AOCC_ROOT_DIR='/cm/shared/apps/spack/cpu/opt/
spack/linux-centos8-zen/gcc-8.3.1/aocc-2.2.0-
hikjjf4adelw56s3namnvyplezykf4fi'
gcc -O3 -mtune=native savage.c -o savagetc.x "-I${AOCC_ROOT_DIR}/include" "-L${AOCC_ROOT_DIR}/lib" -
lamdlibm -lm
icc -O3 -mtune=native savage.c -o savagetc.x "-I${AOCC_ROOT_DIR}/include" "-L${AOCC_ROOT_DIR}/lib" -
lamdlibm -lm
clang -O3 -mtune=native savage.c -o savagetc.x "-I${AOCC_ROOT_DIR}/include" "-L${AOCC_ROOT_DIR}/lib" -
lamdlibm -lm
```

Savage Benchmark: Other Observations & Assumptions

-ffast-math likely simplifies savage expression at compile time

```
gcc -O3 -ffast-math -mtune=native savage.c -o savagec.x -lm
time -p ./savagec.x 1000000000
n = 10000000000
a = 10000000000.465963
5 real 46.98
user 46.82
sys 0.00
```

fixed point at high iteration count is due to 64-bit limit

```
gfortran -ffree-form -ffree-line-length-none -fimplicit-none
-O3 -mtune=native -fdefault-integer-8 -fdefault-real-16
savage.f90 -o savagef.x -lm
time -p ./savagef.x 1000000000
n = 10000000000
a = 999999999.99999985545811435088814878
5 real 3318.56
user 3302.81
sys 0.00
```

Example: Matrix Multiplication & AMD BLIS

PAVE Benchmarks

A collection of linear algebra benchmarks using the **Hilbert matrix**,
a **square matrix** with elements that are **unit fractions** given by

$$H_{ij} = \frac{1}{i+j-1}$$

For example, the Hilbert matrix of dimension 4 is

$$H = \begin{bmatrix} 1 & 1/2 & 1/3 & 1/4 \\ 1/2 & 1/3 & 1/4 & 1/5 \\ 1/3 & 1/4 & 1/5 & 1/6 \\ 1/4 & 1/5 & 1/6 & 1/7 \end{bmatrix}$$

<https://github.com/mkandes/pave>

PAVE Benchmark: HxH

Compute the square of the Hilbert matrix of dimension n .

$$H_{ij}^2 = \sum_{k=1}^n H_{ik} H_{kj}$$

It can be shown that the first element of resultant matrix in the limit of large n approaches the summation of the reciprocals of the squares of the natural numbers. See [Basel problem](#).

$$\sum_{n=1}^{\infty} \frac{1}{n^2} = \frac{\pi^2}{6} \approx 1.644934$$

PAVE Benchmark: HxH - MATMUL

COMPILER	H_{11}^2	AVG (s)	STDEV (s)
gfortran 10.2.0	1.6448730335545843	649.636	26.123
ifort 19.1.1.217	1.64487303355459	4432.478	881.341
aocc/flang 2.2.0	1.644873033554583	1430.356	400.068

```
gfortran -ffree-form -ffree-line-length-none -fimplicit-none  
-O3 -ffast-math -mtune=native -fdefault-integer=8 -  
fdefault-real=8 -fdefault-double=8 hxh-matmul.f90 -o hxh  
-matmul.x

ifort -implicitnone -free -stand none -module build -integer  
-size 64 -real-size 64 -double-size 64 -O3 -fast -mtune=  
native hxh-matmul.f90 -o hxh-matmul.x

5 flang -Mfreeform -Mstandard -fdefault-integer=8 -fdefault-  
real=8 -O3 -ffast-math -mtune=native hxh-matmul.f90 -o  
hxh-matmul.x
```

PAVE Benchmark: HxH - DGEMM

gfortran 10.2.0

BLAS	H_{11}^2	AVG (s)	STDEV (s)
Netlib 3.9.1	1.6448730335545856	3477.714	2101.324
OpenBLAS 0.3.10	1.6448730335545845	250.870	39.788
Intel MKL 2019.1.144	1.6448730335545829	206.712	32.682
AMD BLIS 2.2	1.6448730335545845	244.464	54.536

```
gfortran -ffree-form -ffree-line-length-none -fimplicit-none
          -O3 -ffast-math -mtune=native -fdefault-integer-8 -
          fdefault-real-8 -fdefault-double-8 hxh-dgemm.f90 -o hxh-
          dgemm.x -lopenblas

gfortran -ffree-form -ffree-line-length-none -fimplicit-none
          -O3 -ffast-math -mtune=native -fdefault-integer-8 -
          fdefault-real-8 -fdefault-double-8 hxh-dgemm.f90 -o hxh-
          dgemm.x -lmkl_avx2 -lmkl_core -lmkl_gf_lp64 -
          lmkl_sequential

5 gfortran -ffree-form -ffree-line-length-none -fimplicit-none
          -O3 -ffast-math -mtune=native -fdefault-integer-8 -
          fdefault-real-8 -fdefault-double-8 hxh-dgemm.f90 -o hxh-
          dgemm.x "-L${AOCL_ROOT_DIR}/lib" -lblis -lpthread
```

Conclusions and Summary

- ▶ AMD is developing its own C/C++ and Fortran compiler infrastructure built on the LLVM toolchain.
- ▶ A number of optimized math libraries for high-performance computing are already available and under active development.
- ▶ Where your code will find performance gains is difficult to say without testing, but you should familiarize yourself with the AMD compilers and math libraries to make that assessment
- ▶ Next steps (for me): Learn how to use AMD LibFLAME, then retry with `gpecn1Dtdpbc`.

AMD References

- ▶ AOCC 2.2 User Guide
- ▶ AOCC 3.0 User Guide
- ▶ AOCC 2.3 Clang C/C++ Compiler
- ▶ AOCC 3.0 Clang C/C++ Compiler
- ▶ AOCC 2.3 Flang Fortran Compiler
- ▶ AOCC 3.0 Flang Fortran Compiler
- ▶ AOCL 2.2 User Guide
- ▶ AOCL 3.0 User Guide

AMD References

- ▶ AMD EPYC 7xx2-series Processor Compiler Options Quick Reference Guide
- ▶ <https://developer.amd.com/amd-aocc>
- ▶ <https://developer.amd.com/amd-aocl>
- ▶ <https://developer.amd.com/resources/epyc-resources/epyc-tuning-guides>
- ▶ <https://developer.amd.com/resources/developer-guides-manuals>
- ▶ <https://github.com/amd>

Additional References

- ▶ Partnership for Advanced Computing in Europe
- ▶ PRACE Best Practice Guide - AMD EPYC by Xu Guo

Questions?